
7 Introduction to Recursion

My goal in this chapter is to write a procedure named `downup` that behaves like this:

```
? downup "hello
hello
hell
hel
he
h
he
hel
hell
hello
? downup "goodbye
goodbye
goodby
goodb
good
goo
go
g
go
goo
good
goodb
goodby
goodbye
```

The programming techniques we've used so far in this book don't allow an elegant solution to this problem. We'll use a new technique called *recursion*: writing a procedure that uses *itself* as a subprocedure.

We're going to solve this problem using recursion. It turns out that the idea of recursion is both very powerful—we can solve a *lot* of problems using it—and rather tricky to understand. That's why I'm going to explain recursion several different ways in the coming chapters. Even after you understand one of them, you'll probably find that thinking about recursion from another point of view enriches your ability to use this idea. The explanation in this chapter is based on the *combining method*.

Starting Small

My own favorite way to understand recursion is based on the general problem-solving strategy of solving a complicated problem by starting with a simpler version. To solve the `downup` problem, I'll start by solving this simpler version: write a `downup` procedure that works only for a single-character input word. (You can't get much simpler than that!) Here's my solution:

```
to downup1 :word
  print :word
end
```

```
? downup1 "j
j
```

See how well it works?

Building Up

Of course, `downup1` won't work at all if you give it an input longer than one character. You may not think this was such a big step. But bear with me. Next I'll write a procedure that acts like `downup` when you give it a two-letter input word:

```
to downup2 :word
  print :word
  print butlast :word
  print :word
end
```

```
? downup2 "it
it
i
it
```

We could keep this up for longer and longer input words, but each procedure gets more and more complicated. Here's `downup3`:

```
to downup3 :word
print :word
print butlast :word
print butlast butlast :word
print butlast :word
print :word
end
```

☞ How many `print` instructions would I need to write `downup4` this way? How many would I need for `downup20`?

Luckily there's an easier way. Look at the result of invoking `downup3`:

```
? downup3 "dot
dot
do
d
do
dot
```

The trick is to recognize that the boxed lines are what we'd get by invoking `downup2` with the word `do` as input. So we can find the instructions in `downup3` that print those three lines and replace them with one instruction that calls `downup2`:

```
to downup3 :word
print :word
downup2 butlast :word
print :word
end
```

You might have to think a moment to work out where the `butlast` came from, but consider that we're given the word `dot` and we want the word `do`.

Once we've had this idea, it's easy to extend it to longer words:

```
to downup4 :word
print :word
downup3 butlast :word
print :word
end
```

```
to downup5 :word
  print :word
  downup4 butlast :word
  print :word
end
```

☞ Can you rewrite `downup2` so that it looks like these others?

☞ Before going on, make sure you really understand these procedures by answering these questions: What happens if you use one of these numbered versions of `downup` with an input that is too long? What if the input is too short?

Generalizing the Pattern

We're now in good shape as long as we want to `downup` short words. We can pick the right version of `downup` for the length of the word we have in mind:

```
? downup5 "hello
hello
hell
hel
he
h
he
hel
hell
hello
? downup7 "goodbye
goodbye
goodby
goodb
good
goo
go
g
go
goo
good
goodb
goodby
goodbye
```

Having to count the number of characters in the word is a little unaesthetic, but we could even have the computer do that for us:

```
to downup :word
if equalp count :word 1 [downup1 :word]
if equalp count :word 2 [downup2 :word]
if equalp count :word 3 [downup3 :word]
if equalp count :word 4 [downup4 :word]
if equalp count :word 5 [downup5 :word]
if equalp count :word 6 [downup6 :word]
if equalp count :word 7 [downup7 :word]
end
```

There's only one problem. What if we want to be able to say

```
downup "antidisestablishmentarianism
```

You wouldn't want to have to type in separate versions of `downup` all the way up to `downup28`!

What I hope you're tempted to do is to take advantage of the similarity of all the numbered `downup` procedures by combining them into a single procedure that looks like this:

```
to downup :word
print :word
downup butlast :word
print :word
end
```

(Remember that Logo's `to` command won't let you redefine `downup` if you've already typed in my earlier version with all the `if` instruction lines. Before you can type in the new version, you have to `erase` the old one.)

Compare this version of `downup` with one of the numbered procedures like `downup5`. Do you see that this combined version should work just as well, if all the numbered `downup` procedures are identical except for the numbers in the procedure names? Convince yourself that that makes sense.

☞ Okay, now try it.

What Went Wrong?

You probably saw something like this:

```
? downup "hello
hello
hell
hel
he
h
```


butlast doesn't like as input in downup

There's nothing wrong with the reasoning I used in the last section. If all the numbered `downup` procedures are identical except for the numbers, it should work to replace them all with a single procedure following the same pattern.

The trouble is that the numbered `downup` procedures *aren't* quite all identical. The exception is `downup1`. If it were like the others, it would look like this:

```
to downup1 :word
  print :word
  downup0 butlast :word
  print :word
end
```

Review the way the numbered `downups` work to make sure you understand why `downup1` is different. Here's what happens when you invoke one of the numbered versions:



In this chart, instructions within a particular procedure are indented the same amount. For example, the lines `print "hello` and `downup4 "hell` are part of `downup5`, as is the line `print "hello` at the very end of the chart. The lines in between are indented more because they're part of `downup4` and its subprocedures.

(By the way, the lines in the chart don't show actual instructions in the procedure definitions. Otherwise all the `print` lines would say `print :word` instead of showing actual words. In the chart I've already evaluated the inputs to the commands.)

The point of the chart is that `downup1` has to be special because it marks the end of the “down” part of the problem and the beginning of the “up” part. `downup1` doesn't invoke a lower-numbered `downup` subprocedure because there's no smaller piece of the word to print.

☞ Okay, Logo knows when to stop the “down” part of the program because `downup1` is different from the other procedures. Question: How does Logo know when to stop the “up” part of the program? Why doesn't `downup5`, in this example, have to be written differently from the others?

The Stop Rule

Our attempt to write a completely general `downup` procedure has run into trouble because we have to distinguish two cases: the special case in which the input word contains only one character and the general case for longer input words. We can use `ifelse` to distinguish the two cases:

```
to downup :word
ifelse equalp count :word 1 [downup.one :word] [downup.many :word]
end
```

```
to downup.one :word
print :word
end
```

```
to downup.many :word
print :word
downup butlast :word
print :word
end
```

You'll find that this version of the `downup` program actually works correctly. Subprocedure `downup.one` is exactly like the old `downup1`, while `downup.many` is like the version of `downup` that didn't work.

It's possible to use the same general idea, however—distinguishing the special case of a one-letter word—without having to set up this three-procedure structure. Instead we can take advantage of the fact that `downup.one`'s single instruction is the same as the first instruction of `downup.many`; we can use a single procedure that **stops** early if appropriate.

```
to downup :word
  print :word
  if equalp count :word 1 [stop]
  downup butlast :word
  print :word
end
```

The `if` instruction in this final version of `downup` is called a *stop rule*.

`Downup` illustrates the usual pattern of a recursive procedure. There are three kinds of instructions within its definition: (1) There are the ordinary instructions that carry out the work of the procedure for a particular value of the input, in this case the `print` instructions. (2) There is at least one *recursive call*, an instruction that invokes the same procedure with a smaller input. (3) There is a stop rule, which prevents the recursive invocation when the input is too small.

It's important to understand that the stop rule always comes *before* the recursive call or calls. One of the common mistakes made by programmers who are just learning about recursion is to think this way: "The stop rule *ends* the program, so it belongs at the *end* of the procedure." The right way to think about it is that the purpose of the stop rule is to stop the innermost invocation of the procedure *before* it has a chance to invoke itself recursively, so the stop rule must come *before* the recursive call.

Local Variables

When you're thinking about a recursive procedure, it's especially important to remember that each invocation of a procedure has its own local variables. It's possible to get confused about this because, of course, if a procedure invokes itself as a subprocedure, each invocation uses the same *names* for local variables. For example, each invocation of `downup` has a local variable (its input) named `word`. But each invocation has a *separate* input variable.

It's hard to talk about different invocations in the abstract. So let's look back at the version of the program in which each invocation had a different procedure name: `downup1`, `downup2`, and so on.

If you type the instruction

```
downup5 "hello
```

the procedure `downup5` is invoked, with the word `hello` as its input. `Downup5` has a local variable named `word`, which contains `hello` as its value. The first instruction in `downup5` is

```
print :word
```

Since `:word` is `hello`, this instruction prints `hello`. The next instruction is

```
downup4 butlast :word
```

This instruction invokes procedure `downup4` with the word `hell` (the `butlast` of `hello`) as input. `downup4` has a local variable that is also named `word`. The value of *that* variable is the word `hell`.

At this point there are two separate variables, both named `word`. `Downup5`'s `word` contains `hello`; `downup4`'s `word` contains `hell`. I won't go through all the details of how `downup4` invokes `downup3` and so on. But eventually `downup4` finishes its task, and `downup5` continues with its final instruction, which is

```
print :word
```

Even though different values have been assigned to variables named `word` in the interim, *this* variable named `word` (the one that is local to `downup5`) still has its original value, `hello`. So that's what's printed.

In the recursive version of the program exactly the same thing happens about local variables. It's a little harder to describe, because all the procedure invocations are invocations of the same procedure, `downup`. So I can't say things like "the variable `word` that belongs to `downup4`"; instead, you have to think about "the variable named `word` that belongs to the second invocation of `downup`." But even though there is only one *procedure* involved, there are still five procedure *invocations*, each with its own local variable named `word`.

More Examples

☞ Before I go on to show you another example of a recursive procedure, you might try to write `down` and `up`, which should work like this:

```
? down "hello
hello
hell
hel
he
h
? up "hello
h
he
hel
hell
hello
```

As a start, notice that there are two `print` instructions in `downup` and that one of them does the “down” half and the other does the “up” half. But you’ll find that just eliminating one of the `prints` for `down` and the other for `up` doesn’t *quite* work.

After you’ve finished `down` and `up`, come back here for a discussion of a similar project, which I call `inout`:

```
? inout "hello
hello
  ello
    llo
      lo
        o
        lo
      llo
    ello
  hello
```

At first glance `inout` looks just like `downup`, except that it uses the `butfirst` of its input instead of the `butlast`. `Inout` is somewhat more complicated than `downup`, however, because it has to print spaces before some of the words in order to line up the rightmost letters. `Downup` lined up the leftmost letters, which is easy.

Suppose we start, as we did for `downup`, with a version that only works for single-letter words:

```
to inout1 :word
print :word
end
```

But we can't quite use `inout1` as a subprocedure of `inout2`, as we did in the `downup` problem. Instead we need a different version of `inout1`, which types a space before its input:

```
to inout2 :word
print :word
inout2.1 butfirst :word
print :word
end
```

```
to inout2.1 :word
type "| |" ; a word containing a space
print :word
end
```

`Type` is a command, which requires one input. The input can be any datum. `Type` prints its input, like `print`, but does not move the cursor to a new line afterward. The cursor remains right after the printed datum, so the next `print` or `type` command will continue on the same line.

We need another specific case or two before a general pattern will become apparent. Here is the version for three-letter words:

```
to inout3 :word
print :word
inout3.2 butfirst :word
print :word
end
```

```
to inout3.2 :word
type "| |"
print :word
inout3.1 butfirst :word
type "| |"
print :word
end
```

```

to inout3.1 :word
repeat 2 [type " | "]
print :word
end

```

Convince yourself that each of these procedures types the right number of spaces before its input word.

Here is one final example, the version for four-letter words:

```

to inout4 :word
print :word
inout4.3 butfirst :word
print :word
end

```

```

to inout4.3 :word
type " | |
print :word
inout4.2 butfirst :word
type " | |
print :word
end

```

```

to inout4.2 :word
repeat 2 [type " | "]
print :word
inout4.1 butfirst :word
repeat 2 [type " | "]
print :word
end

```

```

to inout4.1 :word
repeat 3 [type " | "]
print :word
end

```

☞ Try this out and try writing `inout5` along the same lines.

How can we find a common pattern that will combine the elements of all these procedures? It will have to look something like this:

```

to inout :word
repeat something [type " | |]
print :word
if something [stop]
inout butfirst :word
repeat something [type " | |]
print :word
end

```

This is not a finished procedure because we haven't figured out how to fill the blanks. First I should remark that the stop rule is where it is, after the first `print`, because that's how far the innermost procedures (`inout2.1`, `inout3.1`, and `inout4.1`) get. They type some spaces, print the input word, and that's all.

Another thing to remark is that the first input to the `repeat` commands in this general procedure will sometimes be zero, because the outermost procedures (`inout2`, `inout3`, and `inout4`) don't type any spaces at all. Each subprocedure types one more space than its superprocedure. For example, `inout4` types no spaces. Its subprocedure `inout4.3` types one space. `inout4.3`'s subprocedure `inout4.2` types two spaces. Finally, `inout4.2`'s subprocedure `inout4.1` types three spaces.

In order to vary the number of spaces in this way, the solution is to use another input that will have this number as its value. We can call it `spaces`. The procedure will then look like this:

```

to inout :word :spaces
repeat :spaces [type " | |]
print :word
if equalp count :word 1 [stop]
inout (butfirst :word) (:spaces+1)
repeat :spaces [type " | |]
print :word
end

```

```

? inout "hello 0
hello
ello
llo
lo
o
lo
llo
ello
hello

```

Notice that, when we use `inout`, we have to give it a zero as its second input. We could eliminate this annoyance by writing a new `inout` that invokes this one as a subprocedure:

```
to inout :word
  inout.sub :word 0
end

to inout.sub :word :spaces
  repeat :spaces [type "| "]
  print :word
  if equalp count :word 1 [stop]
  inout.sub (butfirst :word) (:spaces+1)
  repeat :spaces [type "| "]
  print :word
end
```

(The easiest way to make this change is to edit `inout` with the Logo editor and change its title line and its recursive call so that its name is `inout.sub`. Then, still in the editor, type in the new superprocedure `inout`. When you leave the editor, both procedures will get their new definitions.)

This program structure, with a short superprocedure and a recursive subprocedure, is very common. The superprocedure's only job is to provide the initial values for some of the subprocedure's inputs, so it's sometimes called an *initialization procedure*. In this program `inout` is an initialization procedure for `inout.sub`.

By the way, the parentheses in the recursive call aren't really needed; I just used them to make it more obvious which input is which.

Other Stop Rules

The examples I've shown so far use this stop rule:

```
if equalp count :word 1 [stop]
```

Perhaps you wrote your `down` procedure the same way:

```
to down :word
  print :word
  if equalp count :word 1 [stop]
  down butlast :word
end
```

Here is another way to write `down`, which has the same effect. But this is a more commonly used style:

```
to down :word
if empty? :word [stop]
print :word
down butlast :word
end
```

This version of `down` has the `stop` rule as its first instruction. After that comes the instructions that carry out the specific work of the procedure, in this case the `print` instruction. The recursive call comes as the last instruction.

A procedure in which the recursive call is the last instruction is called *tail recursive*. We'll have more to say later about the meaning of tail recursion. (Actually, to be precise, I should have said that a *command* in which the recursive call is the last instruction is tail recursive. What constitutes a tail recursive operation is a little trickier, and so far we haven't talked about recursive operations at all.)

Here's another example:

```
to countdown :number
if equal? :number 0 [print "Blastoff! stop]
print :number
countdown :number-1
end
```

```
? countdown 10
10
9
8
7
6
5
4
3
2
1
Blastoff!
```

In this case, instead of a word that gets smaller by `butfirsting` or `butlasting` it, the input is a number from which 1 is subtracted for each recursive invocation. This example

also shows how some special action (the `print "Blastoff!"` instruction) can be taken in the innermost invocation of the procedure.

☞ Here are some ideas for recursive programs you can write. In each case I'll show an example or two of what the program should do. Start with `one.per.line`, a command with one input. If the input is a word, the procedure should print each letter of the word on a separate line. If the input is a list, the procedure should print each member of the list on a separate line:

```
? one.per.line "hello
h
e
l
l
o
? one.per.line [the rain in spain]
the
rain
in
spain
```

(You already know how to do this without recursion, using `foreach` instead. Many, although not all, recursive problems can also be solved using higher order functions. You might enjoy this non-obvious example:

```
to down :word
ignore cascade (count :word) [print ? butlast ?] :word
end
```

While you're learning about recursion, though, don't use higher order functions. Once you're comfortable with both techniques you can choose which to use in a particular situation.)

☞ As an example in which an initialization procedure will be helpful, try `triangle`, a command that takes a word as its single input. It prints the word repeatedly on the same line, as many times as its length. Then it prints a second line with one fewer repetition, and so on until it prints the word just once:


```
? triangle "frog
frog frog frog frog
frog frog frog
frog frog
frog
```

☞ A more ambitious project is **diamond**, which takes as its input a word with an odd number of letters. It displays the word in a diamond pattern, like this:

```
? diamond "program
  g
 ogr
rogra
program
rogra
 ogr
  g
```

(Hint: Write two procedures **diamond.top** and **diamond.bottom** for the growing and shrinking halves of the display. As in **inout**, you'll need an input to count the number of spaces by which to indent each line.) Can you write **diamond** so that it does something sensible for an input word with an even number of letters?

